

?

Bare Demo of IEEEtran.cls for IEEE Journals

Ozlem Erkilic

Pawsey Supercomputing Centre Curtin University
Perth, Australia

Email:ozlem.erkilic@student.curtin.edu.au

Abstract—The best way to understand how to compute a task based on some resources, constraints and goals on advanced computing systems is to work through sample codes. A lot of expert users browse through the Internet to find these examples that they can slightly modify to achieve the results they expect. However, this can be very difficult for some users who are new to this environment as there are always so many of these examples which are not always what they specifically need. This brings out the challenge of having to modify the example to suit their requirements on a particular computing resource, mainly the High Performance Computing systems (HPC) since the modified examples mostly do not work due to HPC systems having different set ups. Therefore, it is hard to tell what the issue is for inexperienced users. Quite often, the users may not be able to tell if the example is broken or if the example is correct, but maybe the code requires some adjustments. This is a very common issue encountered by the users of Pawsey Supercomputing Centre. For this reason, this paper provides solutions to how to run and submit examples on different resources such as supercomputers of Pawsey Supercomputing Centre through the use of a simple tool "getexample". The "getexample" is an effective tool that is developed for the supercomputers of Pawsey that are dependent on a command line interface and aims to help their users learn to how to run code mainly parallel programming examples, submit these examples to different operating systems of Pawsey such as Magnus, Zeus and Zythos. Although, the getexample is limited to these HPC resources, it has the potential to be expanded to other resources and interfaces.

Index Terms—IEEE, IEEEtran, journal, L^AT_EX, paper, template.

I. INTRODUCTION

EVEN though there are many sample codes on search engines for programming purposes, it is usually very rare that these codes work at the first try on different high performance computing systems without modifying them. This is mainly due to each advanced computing system being built up differently from each other to perform specific tasks for the needs of the users. At Pawsey Supercomputing Centre, there are various HPC resources where each of them are set up with slight variations from each other. Therefore, this becomes very challenging for the users, especially the beginners to change the example codes to display a working task on Pawsey's resources as there are many differences between each resources with many constraints. The most common differences between their systems are as follow:

- Contrasting operating systems
- Divergent program environments

- Non-identical compiler options with varying commands
- Various compiler flags and wrappers for varying compiler options
- Module systems and paths that are different
- Disparate module naming conventions
- Various library versions
- Dissimilar personal scratch, group and home directories
- Different scheduling policies

For this reason, when the user runs a sample code found from the Internet for each supercomputers (Magnus, Zeus and Zythos), the code fails to compile as expected and hence, does not run as each individual system has specific operating system and commands set up for them. To rectify these problems, some HPC resources provide their users websites with working examples that aim to assist them how to carry out their tasks step by step. However, sometimes the commands in these examples can be outdated due to the updates in compilers and the operating systems. They also can be challenging to follow and perform on the real systems for the users who are not very experienced with these resources. Although, one may find a well-written bash script to run these systems, may not know how to make this script executable by changing the permissions using the chmod commands. Furthermore, a user may use a sample source code such as a basic MPI code which includes some mpi libraries to perform a particular job on these supercomputers but may fail when compiled due to these libraries being no longer valid or upgraded to a different version. Or one may complete their task, but may not know in which filesystem to store their results, for example in scratch or their group because some supercomputers have policies on how long to keep the data on certain filesystems such as scratch. If they are not moved from that specific file system within a certain amount of time, the results get deleted and thus, the user loses their work. This is a major difficulty for the users, especially researchers and scientist since it takes very long time to collect their data and vital for their researches. It is also very crucial that these are stored correctly within these sources.

In order to minimise these problems, the getexample was developed to supply examples which do not require further editing or modification and works when the executable is run. Thus, it aims to teach the users of Pawsey Supercomputing Centre how to run and submit tasks on the supercomputers without encountering issues. The rest of this paper explains

how the getexample tool works.




Fig. 1. elastic data set 1




Fig. 2. elastic data set 2

II. OVERVIEW OF GETEXAMPLE

The main aim of the "getexample" is to provide easy access to the examples for the users. For this reason, the getexample tool is expected to have the following design specifications:

- It should be easy to use for everyone including beginners. The user should be able to download the examples and run them with one command.
- It should provide practical examples which can be modified or updated by other users for their own work. Therefore, it encourages the users to learn how to use the resources given to them effectively.
- The example should be clear and completely detailed with steps to help the users understand and apply it.

The getexample tool is designed in a way such that when "getexample" command is typed from any directory on the resources of Pawsey Centre, it lists all the examples within the getexample library. As the getexample displays many examples which aim to perform various tasks for each individual resource at Pawsey, it ensures that when the user logs in from a specific supercomputer such as Magnus, it only lists the examples provided for Magnus rather than showing all of the examples within the library. To obtain a copy of any example, the user simply types "getexample" followed by the name of the example as shown below:

```
getexample <name of the example>
```

This creates a new directory with same name as the example requested in wherever the getexample tool is accessed from and it downloads all the files of the example to the new directory created. The new directory then can be accessed by the user simply typing:

```
cd <name of the example>
```

III. CREATING EXAMPLES ON GETEXAMPLE

The examples included in the getexample library were created by internship students at Pawsey with their supervisor while the source codes within the getexample were obtained from some wiki pages and websites and were acknowledged in the examples. Some of the examples were also created due to the needs of research students for their internship projects at Pawsey. The examples obtained in the getexample are mainly for teaching the users how to work with parallel programming including MPI, OpenMP and hybrid jobs such as OpenMP/MPI which utilise basic source codes similar to "Hello world" program and submit these jobs to different supercomputers such as Magnus, Zeus and Zythos using different program environments and compiler modules.

Each individual example on the getexample occupies a directory that is reserved for them and each example consists of three files listed below:

- **SLURM (Simple Linux Utility for Resource Management)** : This allows the users to submit batch jobs to the supercomputers, check on their status and cancel them if needed. It contains the necessary detail about the name of the executable, the results directory, the name of the output file and the jobID. It also allows the users to edit how many nodes the code requires to run on the HPC systems, the duration of the task that it takes, which partition to be used and their account name. The SLURM initially creates a scratch directory for the example to run in and the results are outputted to a log file. Then, it creates a group directory in which the results directory is located for that example. Once the output file is completed within the scratch, the output file is then moved to the results directory located in the group directory and the scratch directory then gets removed.
- **Source Code** : This is usually a source code in c or Fortran and taken from wiki pages to run the example.
- **README** : This file is an executable bash script which can be read and run by the users. It provides details about what the source code does, how to compile the source code depending on the program environment such as Cray, Intel, GNU or PGI, what can be modified in the SLURM directives, and a set of instruction on how to submit the SLURM to the supercomputers including which specific commands to use for particular supercomputers. It can be executed by simply typing ./README which then compiles the source code and submits the batch job to the chosen supercomputer.

For the examples in the getexample tool to be user friendly, helpful and efficient for working with HPC resources, there are many design considerations to be held as listed below:

- Before introducing the examples to the users, it should be made sure that the examples run without encountering any errors. For example, they should be suitable for the current operating systems with the use of correct commands for the different compiler modules such as Intel, GNU, and PGI or different program environments such as Cray.
- All the files to run the example should be included with the example.
- To minimise confusion on how to perform the example on the supercomputers, the example should be executed with a single command. For example, if a simple shell script is used, it should be run as `./README`.
- The examples should have enough instructions about what each command does and which parts of the SLURM and the README file can be modified so that all the files and the scripts should be able to be edited easily by the users for their preferences.
- It must be ensured that the README and SLURM have enough information about how to change certain parts of these files so that the users can run the example how they wish to on the supercomputers. These instructions should be understandable by the new users who are not very experienced with these systems. For example, if the user wishes to use more nodes on the supercomputers, one should be able to know where to change it from.

Even though, the examples in the getexample tool are designed in a such way that once they are downloaded, the user should be able to run them without having to modify them. However, this is not always the case as some of the supercomputers in Pawsey Supercomputing Centre such as Zythos due to different set ups require the account name which is customized for each user and without the correct account name in the SLURM, the code fails to run. Therefore, the examples in the getexample tool ensures that the users are informed on how to change their account name located in the SLURM file. Furthermore, they assist the users on how to change number of nodes used within the supercomputers and increase or reduce the number of cores used

A. Magnus

Magnus is described as a Cray X40 supercomputer that consists of many nodes that are bound by a high speed network. For the compute nodes, each one of them has 2 sockets and each of these has 12 cores. Magnus is also specified to have 24 cores per node and in total these sum up to 35,712 cores across the 1488 nodes. On Magnus, jobs run on the back-end of the system with the help of SLURM and ALPS (the Cray Application Level Placement Scheduler). A batch job is submitted to the queue system on the front-end from the sbatch command. When it runs, it executes the launch command aprun on the MOM nodes which are the login nodes of Magnus. The aprun keeps running on these login nodes until the application gets completed. Once aprun finishes,

the SLURM job also completes. As mentioned previously, the focus of the getexample tool was not only to provide working examples to the users but also assist them on how to run jobs on their scratch directories. Therefore, whenever the batch job was submitted to Magnus, it was ran on the scratch directory and once the job was completed, the results were carried to the group directory. In the end, the scratch was removed. The examples used for Magnus were mainly parallel programming examples such as MPI, OpenMP and hybrid codes which a combination of OpenMP and MPI tasks and some applications such as lammmps and gromacs. The source codes used for these examples were written in c or Fortran and were basic Hello world codes with the exception of the application codes. Each example was displayed in different environments on Magnus such as GNU, Intel and more importantly the default environment, Cray with the compiler options of ftn, mpif90 and cc. When using these environments, it was paramount to load the actual programming environment before compiling the source codes and running them as each environment has specific compiler commands with distinct wrappers. These are explain more in detail in the following sections.

1) *MPI Examples:* MPI is known as a Message Passing Interface application that is a communication model for moving data between processors. For MPI examples, Fortran and c codes were mainly used and ran on different environments including Cray, GNU and Intel. On Magnus, the default program environment is Cray. Therefore, it was necessary to load each program environment that was intended to work on by doing a:

```
module swap PrgEnv-cray PrgEnv-<name of environment>
```

It is also very important to notice that if the right program environment is not loaded, the code will fail as each environment have different commands for compilers. However, on Magnus, swapping from one environment to the other might be very challenging for the new users. As the getexample tool is desired to be as automated as possible to minimise failures on tasks, it was made sure that after running each job on Intel and GNU environments, the environment was set back to the default environment, Cray. This prevents the users to manually list the module every time they run something on Magnus and modify the codes to swap from one module to the other.

The very first MPI example performed on Magnus ran on 2 nodes with a total of 48 tasks with a basic "Hello world" source code in both c and Fortran. This example consisted of three files as mentioned previously which were README, SLURM and the source codes. The SLURM script included information about the choice in the number of nodes, the partition, the duration of the time it takes to run the job, where to run the task such as on scratch and where to store the results, for example the group. It also specified the name of the executable, the results directory and the output file.

For Magnus, the classified partition is the debugq. Since, there were 2 nodes used, the SLURM directives were given as:

```
#!/bin/bash -l
#SBATCH --job-name=GE-hostname
#SBATCH --partition=debugq
#SBATCH --nodes=2
#SBATCH --time=00:05:00
#SBATCH --export=NONE
```

The generic variables such as the executable, scratch, the results directory and the output file were defined as shown below:

```
EXECUTABLE=hello_mpi_cray
SCRATCH=$MYSCRATCH/run_hostname/$SLURM_JOBID
RESULTS=$MYGROUP/mpifortran_cray_results/$SLURM_JOBID

OUTPUT=mpifortran_cray.log
```

To launch MPI job with fully occupied 2 nodes, the aprun command was used as expressed below:

```
aprun -n 48 -N 24 ./EXECUTABLE >> ${OUTPUT}
```

where -n defines the total number of MPI tasks while -N specifies the number of MPI tasks per node as there are 2 nodes.

In the README file, the correct program environment should be included for the codes to run well. For compiling on Cray, there was no need to load this environment as the default module was Cray. However, to run on GNU and Intel, the program environment was changed from Cray to GNU or Intel as shown below:

```
module swap PrgEnv-cray PrgEnv-gnu
module swap PrgEnv-cray PrgEnv-intel
```

However, the compiler commands do not change for Cray, GNU and Intel for MPI codes on Magnus, whereas they have distinct differences for other parallel programming tasks such as OpenMP.

To compile the "Hello world" MPI Fortran code, hello_mpi.f90 on Cray, GNU and Intel environments:

```
ftn -O2 hello_mpi.f90 -o hello_mpi_cray
```

-O2 is the optimization method, while -o placed after the source code directs to an executable called hello_mpi_cray which was previously defined in the SLURM and chosen by the user as name. For GNU and Intel, this was called hello_mpi_gnu and hello_mpi_intel.

When the source code in use was the c code, the SLURM was not different from that of Fortran. However, the README did change. To compile the hello_mpi.c code, the following command was used:

```
cc -O2 hello_mpi.c -o hello_mpi_gnu
```

Once the codes were compiled, the SLURM was then submitted to Magnus by:

```
sbatch hello_mpi_cray.slurm
```

As mentioned earlier, to prevent the users from manually listing the modules to see which modules are loaded and from

which module to swap, at the end of the SLURM and the README, the program environment was always set back to the default, Cray by:

```
module swap PrgEnv-<name of environment> PrgEnv-cray
```

Another example on MPI was to run the same sources but on partially occupied nodes which means that instead of having 24 tasks per node, each node could only have 12 tasks. When running this example on Magnus, the README remained unchanged, but there were minor changes to the SLURM.

The number of OpenMP threads was set to 1 to prevent from inadvertent OpenMP threading and the aprun command was changed to:

```
aprun -n 24 -N 12 -S 6 ./EXECUTABLE >> ${OUTPUT}
```

-n defines the total number of MPI tasks, -N specifies the number tasks per node while -S defines 6 MPI tasks per socket.

2) *OMP Examples:* OpenMP (Open Multi-Processing) is known as a common shared memory model in which the threads work in parallel and access all shared memory. Therefore, it can be thought of one task because of their restrictions to only one node. Unlike the MPI examples, the compiler commands and the wrappers do change for different environments on Magnus with the OpenMP tasks. Thus, it is necessary to swap to the right program environment corresponding to the compiler.

As mentioned previously, the OpenMP jobs use only one node and this node can be fully occupied node or a node occupying a single NUMA region. When it was fully occupied, the slurm script had a choice of only one node with the debugq partition as explained.

```
#SBATCH --partition=debugq
#SBATCH --nodes=1
#SBATCH --time=00:05:00
#SBATCH --export=NONE
```

The initiation of the job by aprun is done by specifying that there is one task and 24 threads. The number of threads is also called a "depth" and hence, the number of threads, OMP_NUM_THREADS was set to 24 giving an expression of:

```
export OMP_NUM_THREADS=24
aprun -n 1 -d 24 ./EXECUTABLE >> ${OUTPUT}
```

For a node occupying a single NUMA region, the code are more efficient when threads are limited to one NUMA region containing 12 cores. The expression (-d) 12 ensures that threads are bound correctly and this can be clearly done by specifying this via -cc which specifies the cores used. This is all illustrated below as:

```
export OMP_NUM_THREADS=12
aprun -n 1 -d 12 -cc 0-11 ./EXECUTABLE >> ${OUTPUT}
```

The source codes used for OpenMP jobs were omp_hello.f and omp_hello.c. To run these source codes respectively on